

HW2

Nathan Krieger

2026-01-30

Problem 1

(a)

$$\hat{\beta}x_i = \hat{\alpha}z_i \quad \hat{\beta} = c\hat{\alpha}$$

$$\hat{\alpha}_j = \frac{\hat{\beta}_j}{c}$$

(b)

(c)

(d)

I would say this is incorrect because the student forgot the ϵ .

(e)

False. You can make the training model whatever size you want.

(f)

Unbiased means that if you did the experiment a large amount of times and calculate $\hat{\beta}$ for each of those times, the average of all the estimates would exactly equal β .

Problem 2

```
#install.packages("ISLR2") #you only need to do this one time.  
library(ISLR2) #you will need to do this every time you open a new R session.
```

```
## Warning: package 'ISLR2' was built under R version 4.5.2
```

(a)

```
head(Boston)
```

```
##      crim zn indus chas   nox     rm    age    dis rad tax ptratio lstat medv  
## 1 0.00632 18  2.31    0 0.538 6.575 65.2 4.0900    1 296 15.3 4.98 24.0  
## 2 0.02731  0  7.07    0 0.469 6.421 78.9 4.9671    2 242 17.8 9.14 21.6  
## 3 0.02729  0  7.07    0 0.469 7.185 61.1 4.9671    2 242 17.8 4.03 34.7  
## 4 0.03237  0  2.18    0 0.458 6.998 45.8 6.0622    3 222 18.7 2.94 33.4  
## 5 0.06905  0  2.18    0 0.458 7.147 54.2 6.0622    3 222 18.7 5.33 36.2  
## 6 0.02985  0  2.18    0 0.458 6.430 58.7 6.0622    3 222 18.7 5.21 28.7
```

```

set.seed(1)

n <- nrow(Boston)

train_index <- sample(1:n, floor(n / 2), replace = FALSE)

train_boston <- Boston[train_index, ]
test_boston <- Boston[-train_index, ]

m1 <- lm(crim ~ ., data = train_boston)

summary(m1)

##
## Call:
## lm(formula = crim ~ ., data = train_boston)
##
## Residuals:
##    Min      1Q  Median      3Q     Max 
## -10.574  -2.723  -0.566   1.351  57.279 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 18.665277 10.693890  1.745  0.08219 .
## zn          0.046105  0.028198  1.635  0.10336  
## indus       -0.127032  0.116665 -1.089  0.27730  
## chas        -0.916885  1.769171 -0.518  0.60476  
## nox         -11.606805  7.924234 -1.465  0.14431  
## rm          0.738859  0.913675  0.809  0.41951  
## age         -0.010585  0.026291 -0.403  0.68761  
## dis         -1.184115  0.427288 -2.771  0.00602 ** 
## rad          0.671788  0.130702  5.140  5.7e-07 *** 
## tax          -0.004607  0.007552 -0.610  0.54237  
## ptratio      -0.515160  0.284824 -1.809  0.07175 .  
## lstat        0.296310  0.114591  2.586  0.01031 *  
## medv        -0.249594  0.097588 -2.558  0.01115 * 
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.693 on 240 degrees of freedom
## Multiple R-squared:  0.5021, Adjusted R-squared:  0.4772 
## F-statistic: 20.17 on 12 and 240 DF,  p-value: < 2.2e-16

trainMSE <- sum(m1$residuals^2)/nrow(train_boston)

trainMSE

## [1] 42.49345

test_predict <- predict(m1, newdata = test_boston)

test_MSE <- mean((test_predict - test_boston$crim)^2)

test_MSE

```

```

## [1] 41.19923

(b)

set.seed(1)

n <- nrow(Boston)

train_index <- sample(1:n, floor(n / 2), replace = FALSE)

train_boston <- Boston[train_index, ]
test_boston <- Boston[-train_index, ]

m1 <- lm(crim ~ zn + indus + nox + dis + rad + ptratio + medv, data = train_boston)

summary(m1)

##
## Call:
## lm(formula = crim ~ zn + indus + nox + dis + rad + ptratio +
##     medv, data = train_boston)
##
## Residuals:
##    Min      1Q  Median      3Q     Max 
## -9.145 -2.542 -0.645  1.175 57.028 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 29.29068   8.84292   3.312 0.001065 ** 
## zn          0.04940   0.02779   1.777 0.076733 .    
## indus       -0.13410   0.10956  -1.224 0.222122    
## nox        -12.25307   7.66941  -1.598 0.111408    
## dis         -1.36389   0.39946  -3.414 0.000748 ***  
## rad          0.63065   0.07015   8.990 < 2e-16 ***  
## ptratio     -0.57530   0.28209  -2.039 0.042483 *   
## medv        -0.35150   0.06458  -5.443 1.27e-07 *** 
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.739 on 245 degrees of freedom
## Multiple R-squared:  0.4847, Adjusted R-squared:  0.47 
## F-statistic: 32.93 on 7 and 245 DF,  p-value: < 2.2e-16

trainMSE <- sum(m1$residuals^2)/nrow(train_boston)

trainMSE

## [1] 43.97466

test_predict <- predict(m1, newdata = test_boston)

test_MSE <- mean((test_predict - test_boston$crim)^2)

test_MSE

```

```
## [1] 39.62763
```

The training MSE collected from part A was slightly smaller than the one from part B. The test MSE collected from part A was slightly larger than the one from part B.

(c)

I expected part b to have a larger MSE because it is being trained on less variables so it's not likely to be as good as part a that was trained on every variable.

(d)

Problem 3

(a)

$$\beta_0 = 2 \quad \beta_1 = 3 \quad \beta_2 = 5$$

(b)

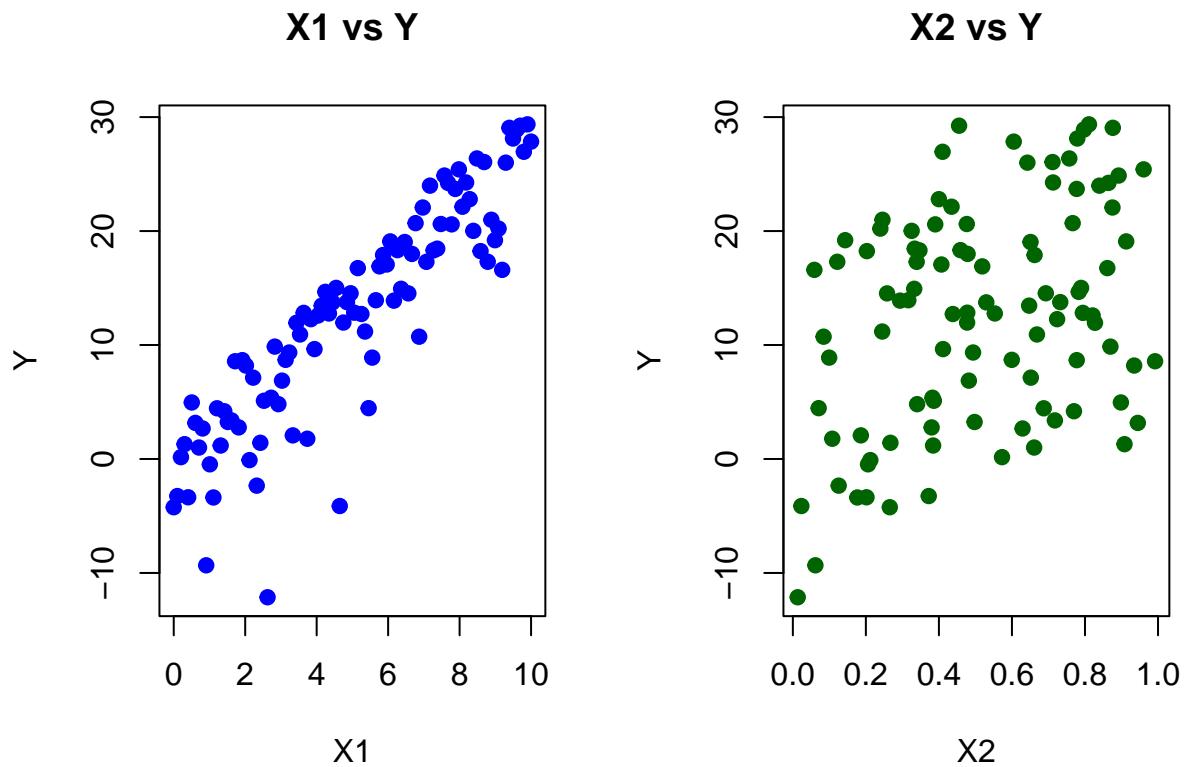
```
set.seed(1) # For reproducibility
X1 = seq(0, 10, length.out = 100)
X2 = runif(100)
epsilon = rnorm(100, mean = 0, sd = 1)

Y = 2 + 3*X1 + 5*log(X2) + epsilon
Y
```

	[1]	-4.23243317	-3.25163817	0.16155277	1.29831740	-3.36017219
##	[6]	4.95979463	3.16639013	1.00554051	2.67674841	-9.32815238
##	[11]	-0.46809191	-3.37647122	4.44916434	1.18318278	4.19129737
##	[16]	3.24545027	3.38444034	8.57643582	2.77034153	8.67147697
##	[21]	8.19849529	-0.09879487	7.13633701	-2.33945369	1.42069144
##	[26]	5.10909193	-12.13061519	5.37632558	9.86110241	4.80943921
##	[31]	6.87401554	8.70101323	9.34431309	2.07223625	11.94948013
##	[36]	10.92516778	12.82034189	1.77720644	12.26835426	9.64280773
##	[41]	12.59220416	13.45553046	14.66413305	12.76885860	13.74312926
##	[46]	15.01216225	-4.12701656	11.97037589	13.76311067	14.53952062
##	[51]	12.83644531	16.74957378	12.72008067	11.18201057	4.46102116
##	[56]	8.89426498	13.93063644	16.90011981	17.89753266	17.06416690
##	[61]	19.09030520	13.89557628	18.32735160	14.93305155	19.03933549
##	[66]	14.53050849	17.99498458	20.69307892	10.73023226	22.06593930
##	[71]	17.29846356	23.98309105	18.30688716	18.45521373	20.61605275
##	[76]	24.86960493	24.22779022	20.58752240	23.69519179	25.41423139
##	[81]	22.13662294	24.26178623	22.79845699	20.01888537	26.36971877
##	[86]	18.24079376	26.05506809	17.30404313	20.99204629	19.19887827
##	[91]	20.21505958	16.59568737	26.00022507	29.05787821	28.11965954
##	[96]	28.90449384	29.24380140	26.95436970	29.36243300	27.84620901

(c)

```
par(mfrow=c(1,2))
plot(X1, Y, main="X1 vs Y", col="blue", pch=19)
plot(X2, Y, main="X2 vs Y", col="darkgreen", pch=19)
```



(d)

```
n_sim <- 10000
beta1_hats <- rep(NA, n_sim)

for(i in 1:n_sim) {
  # We generate new noise each time to simulate a new sample
  eps_sim <- rnorm(100, 0, 1)
  Y_sim <- 2 + 3*X1 + 5*log(X2) + eps_sim

  # Fit the model (using log(X2) since that is the true relationship)
  fit <- lm(Y_sim ~ X1 + log(X2))
  beta1_hats[i] <- coef(fit)[2]
}

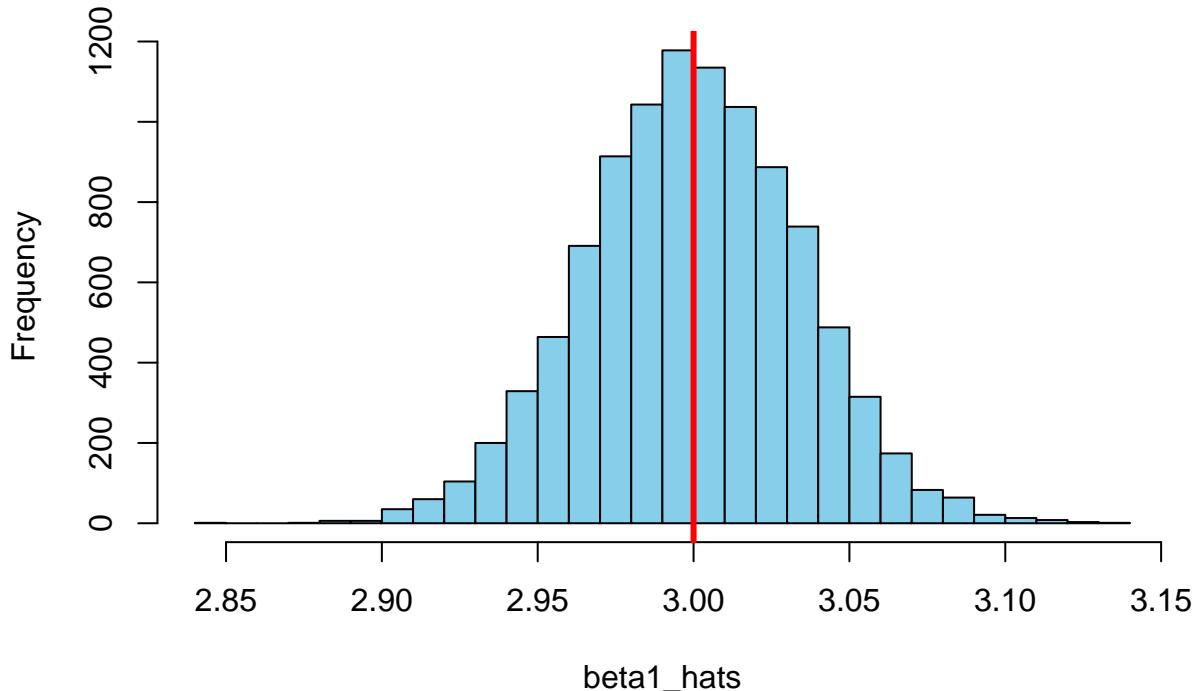
mean(beta1_hats)

## [1] 2.999792
```

(e)

```
hist(beta1_hats, breaks=30, col="skyblue", main="Sampling Distribution of Beta1")
abline(v = 3, col = "red", lwd = 3) # The true Beta1
```

Sampling Distribution of Beta1



```
## (f)
beta2_hats <- rep(NA, n_sim)

for(i in 1:n_sim) {
  eps_sim <- rnorm(100, 0, 1)
  Y_sim <- 2 + 3*X1 + 5*log(X2) + eps_sim
  fit <- lm(Y_sim ~ X1 + log(X2))
  beta2_hats[i] <- coef(fit)[3]
}

mean(beta2_hats)

## [1] 4.999968

(g)

hist(beta2_hats, breaks=30, col="lightgreen", main="Sampling Distribution of Beta2")
abline(v = 5, col = "red", lwd = 3) # The true Beta2
```

Sampling Distribution of Beta2

