

HW2

Nathan Krieger

2026-01-30

Problem 1

(a)

$$\hat{\beta}x_i = \hat{\alpha}z_i$$

$$\hat{\beta} = c\hat{\alpha}$$

$$\hat{\alpha}_j = \frac{\hat{\beta}_j}{c} \text{ for } j = 1 \text{ to } p$$

(b)

No it is not necessary. The least squares estimates automatically adjust the scale to compensate for the units. The fit of the model will be the same.

(c)

Linearity: The relationship between the independent variables and the dependent variable follows a straight-line pattern.

Independence: The data points are collected independently so one person's data doesn't influence another.

Homoscedasticity: The spread or scatter of the errors is the same across all levels of predictors.

Normality: For any value of X, the errors should follow a bell-shaped curve.

(d)

I would say this is incorrect because the student forgot the ϵ .

(e)

False. You can make the training model whatever size you want. The training model is usually smaller but this is not required.

(f)

Unbiased means that on average the model will get the right answer.

If you were to take thousands of different random samples from the same population and calculate $\hat{\beta}$ for each one, the average of all those different estimates would equal the true population value β .

Problem 2

```
#install.packages("ISLR2") #you only need to do this one time.  
library(ISLR2) #you will need to do this every time you open a new R session.
```

(a)

```
head(Boston)

##      crim zn indus chas   nox     rm    age     dis rad tax ptratio lstat medv
## 1 0.00632 18  2.31     0 0.538 6.575 65.2 4.0900    1 296 15.3 4.98 24.0
## 2 0.02731  0  7.07     0 0.469 6.421 78.9 4.9671    2 242 17.8 9.14 21.6
## 3 0.02729  0  7.07     0 0.469 7.185 61.1 4.9671    2 242 17.8 4.03 34.7
## 4 0.03237  0  2.18     0 0.458 6.998 45.8 6.0622    3 222 18.7 2.94 33.4
## 5 0.06905  0  2.18     0 0.458 7.147 54.2 6.0622    3 222 18.7 5.33 36.2
## 6 0.02985  0  2.18     0 0.458 6.430 58.7 6.0622    3 222 18.7 5.21 28.7

set.seed(1)

n <- nrow(Boston)

train_index <- sample(1:n, floor(n / 2), replace = FALSE)

train_boston <- Boston[train_index, ]
test_boston <- Boston[-train_index, ]

m1 <- lm(crim ~ ., data = train_boston)

summary(m1)

##
## Call:
## lm(formula = crim ~ ., data = train_boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -10.574  -2.723  -0.566   1.351  57.279 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 18.665277 10.693890  1.745  0.08219 .
## zn          0.046105  0.028198  1.635  0.10336  
## indus      -0.127032  0.116665 -1.089  0.27730  
## chas        -0.916885  1.769171 -0.518  0.60476  
## nox       -11.606805  7.924234 -1.465  0.14431  
## rm          0.738859  0.913675  0.809  0.41951  
## age        -0.010585  0.026291 -0.403  0.68761  
## dis        -1.184115  0.427288 -2.771  0.00602 ** 
## rad         0.671788  0.130702  5.140  5.7e-07 *** 
## tax        -0.004607  0.007552 -0.610  0.54237  
## ptratio     -0.515160  0.284824 -1.809  0.07175  
## lstat       0.296310  0.114591  2.586  0.01031 *  
## medv      -0.249594  0.097588 -2.558  0.01115 * 
```

```

## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.693 on 240 degrees of freedom
## Multiple R-squared:  0.5021, Adjusted R-squared:  0.4772
## F-statistic: 20.17 on 12 and 240 DF,  p-value: < 2.2e-16
trainMSE <- sum(m1$residuals^2)/nrow(train_boston)

trainMSE

## [1] 42.49345

test_predict <- predict(m1, newdata = test_boston)

test_MSE <- mean((test_predict - test_boston$crim)^2)

test_MSE

## [1] 41.19923

```

(b)

```

m1 <- lm(crim ~ zn + indus + nox + dis + rad + ptratio + medv, data = train_boston)

summary(m1)

##
## Call:
## lm(formula = crim ~ zn + indus + nox + dis + rad + ptratio +
##     medv, data = train_boston)
##
## Residuals:
##    Min      1Q  Median      3Q     Max
## -9.145 -2.542 -0.645  1.175 57.028
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 29.29068   8.84292   3.312 0.001065 **
## zn          0.04940   0.02779   1.777 0.076733 .
## indus       -0.13410   0.10956  -1.224 0.222122
## nox        -12.25307   7.66941  -1.598 0.111408
## dis         -1.36389   0.39946  -3.414 0.000748 ***
## rad          0.63065   0.07015   8.990 < 2e-16 ***
## ptratio     -0.57530   0.28209  -2.039 0.042483 *
## medv        -0.35150   0.06458  -5.443 1.27e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.739 on 245 degrees of freedom
## Multiple R-squared:  0.4847, Adjusted R-squared:  0.47
## F-statistic: 32.93 on 7 and 245 DF,  p-value: < 2.2e-16
trainMSE <- sum(m1$residuals^2)/nrow(train_boston)

trainMSE

```

```

## [1] 43.97466
test_predict <- predict(m1, newdata = test_boston)

test_MSE <- mean((test_predict - test_boston$crim)^2)

test_MSE

## [1] 39.62763

```

The training MSE collected from part A was slightly smaller than the one from part B. The test MSE collected from part A was slightly larger than the one from part B.

(c)

I expected part b to have a larger MSE because it is being trained on less variables so it's not likely to be as good as part a that was trained on every variable.

(d)

I expected the test MSE for part b to be larger than part a because removing predictors usually increases bias. However, if some predictors in part a were just noise, part b could potentially have a smaller test MSE. In my case, part b was slightly better, this means that some variables in part a weren't helpful for generalization.

Problem 3

(a)

$$\beta_0 = 2 \quad \beta_1 = 3 \quad \beta_2 = 5$$

(b)

```

set.seed(1) # For reproducibility
X1 = seq(0, 10, length.out = 100)
X2 = runif(100)
epsilon = rnorm(100, mean = 0, sd = 1)

Y = 2 + 3*X1 + 5*log(X2) + epsilon
Y

##      [1] -4.23243317 -3.25163817  0.16155277  1.29831740 -3.36017219
##      [6]  4.95979463  3.16639013  1.00554051  2.67674841 -9.32815238
##     [11] -0.46809191 -3.37647122  4.44916434  1.18318278  4.19129737
##     [16]  3.24545027  3.38444034  8.57643582  2.77034153  8.67147697
##     [21]  8.19849529 -0.09879487  7.13633701 -2.33945369  1.42069144
##     [26]  5.10909193 -12.13061519  5.37632558  9.86110241  4.80943921
##     [31]  6.87401554  8.70101323  9.34431309  2.07223625 11.94948013
##     [36] 10.92516778 12.82034189  1.77720644 12.26835426  9.64280773
##     [41] 12.59220416 13.45553046 14.66413305 12.76885860 13.74312926
##     [46] 15.01216225 -4.12701656 11.97037589 13.76311067 14.53952062
##     [51] 12.83644531 16.74957378 12.72008067 11.18201057  4.46102116
##     [56]  8.89426498 13.93063644 16.90011981 17.89753266 17.06416690
##     [61] 19.09030520 13.89557628 18.32735160 14.93305155 19.03933549
##     [66] 14.53050849 17.99498458 20.69307892 10.73023226 22.06593930
##     [71] 17.29846356 23.98309105 18.30688716 18.45521373 20.61605275

```

```

## [76] 24.86960493 24.22779022 20.58752240 23.69519179 25.41423139
## [81] 22.13662294 24.26178623 22.79845699 20.01888537 26.36971877
## [86] 18.24079376 26.05506809 17.30404313 20.99204629 19.19887827
## [91] 20.21505958 16.59568737 26.00022507 29.05787821 28.11965954
## [96] 28.90449384 29.24380140 26.95436970 29.36243300 27.84620901


```

(c)

```

par(mfrow=c(1,2))
plot(X1, Y, main="X1 vs Y", col="blue", pch=19)
plot(X2, Y, main="X2 vs Y", col="darkgreen", pch=19)

```



X1 and Y clearly have a positive linear relationship. Also, the graph between X2 and Y shows a much more noisy relationship but slightly resembles a logarithmic relationship.

(d)

```

n_sim <- 10000
beta1_hats <- rep(NA, n_sim)

for(i in 1:n_sim) {
  # We generate new noise each time to simulate a new sample
  eps_sim <- rnorm(100, 0, 1)
  Y_sim <- 2 + 3*X1 + 5*log(X2) + eps_sim

  # Fit the model (using log(X2) since that is the true relationship)
}

```

```

    fit <- lm(Y_sim ~ X1 + log(X2))
    beta1_hats[i] <- coef(fit)[2]
}


mean(beta1_hats)

## [1] 2.999792

(e)

hist(beta1_hats, breaks=30, col="skyblue", main="Sampling Distribution of Beta1")
abline(v = 3, col = "red", lwd = 3) # The true Beta1

```



```

## (f)

beta2_hats <- rep(NA, n_sim)

for(i in 1:n_sim) {
  eps_sim <- rnorm(100, 0, 1)
  Y_sim <- 2 + 3*X1 + 5*log(X2) + eps_sim
  fit <- lm(Y_sim ~ X1 + log(X2))
  beta2_hats[i] <- coef(fit)[3]
}

mean(beta2_hats)

## [1] 4.999968

```

(g)

```
hist(beta2_hats, breaks=30, col="lightgreen", main="Sampling Distribution of Beta2")
abline(v = 5, col = "red", lwd = 3) # The true Beta2
```

Sampling Distribution of Beta2

